


GPU-backed HF



CPU-backed llama.cpp

