



GPU-backed HF



CPU-backed llama.cpp



E0-REF

E1-INT8

E2-Q8

E3-Q5

E4-Q4

Falcon-H1-Tiny-90M

Falcon-H1-Tiny-R-0.6B

Falcon-H1-1.5B-Deep

Qwen2.5-1.5B

Llama-3.2-1B

Phi-3.5-mini

SmolLM2-360M

TinyLlama-1.1B-Chat

Falcon-H1-0.5B

Falcon-H1-1.5B

Qwen2.5-0.5B

Gemma-3-1B

Llama-3.2-3B

Phi-4-mini

SmolLM2-1.7B

InternLM2.5-1.8B-Chat