


Decode throughput (tokens/s, median)

GPU-backed HF



CPU-backed llama.cpp

