

Datasets

SQuAD

HaluEval

TruthfulQA

Natural Question

Safety

MTBench

Models

InternLM2.5-1.8B

Phi 3.5 -Mini
Phi 4 - Mini

Gemma-3-1B

SmolLM2
362M, 1.7B

TinyLlama-1.1B

Qwen2.5
0.5B, 1.5B

Falcon-H1
92M, 0.5B, 0.6B,
1.5B, 1.5B-deep

Llama 3.2
1B, 3B

P0-Expert

P1-Casual

P2-Naive

P3-Broken

Prompt Level

Performance Metric

Exact Match

Hallucination

Accuracy

Latency

Throughput

Safety Refusal

Scheduler

Edge Setting

E0-REF

E1-INT8

E2-Q8

E3-Q5

E4-Q4

HF Transformers

llama.cpp

GPU Inference

CPU Inference