


MT-Bench by Edge Setting per Model Family (P0 — Expert prompt)
E0-REF (GPU BF16) → E4-Q4 (CPU 4-bit) · legend shows Δ vs E0-REF


Falcon-H1

reasoning




Gemma (1B)

reasoning




Llama (1B + 3B)

reasoning




Qwen (0.5B + 1.5B)

reasoning



SmolLM2 (360M + 1.7B)

reasoning



90°

135°

45°

1.0

0.8

0.6

0.4

0.2

0°

180°

225°

270°

315°